Assignment 2: Policy Gradient

Andrew ID: Write your Andrew ID here.
Collaborators: Write the Andrew IDs of your collaborators here (if any).
NOTE: Please do **NOT** change the sizes of the answer blocks or plots.

5 Small-Scale Experiments

5.1 Experiment 1 (Cartpole) – [25 points total]

5.1.1 Configurations

Q5.1.1

```
python rob831/scripts/run_hw2.py --env_name CartPole-v0 -n 100 -b 1000 \
    -dsa --exp_name q1_sb_no_rtg_dsa

python rob831/scripts/run_hw2.py --env_name CartPole-v0 -n 100 -b 1000 \
    -rtg -dsa --exp_name q1_sb_rtg_dsa

python rob831/scripts/run_hw2.py --env_name CartPole-v0 -n 100 -b 1000 \
    -rtg --exp_name q1_sb_rtg_na

python rob831/scripts/run_hw2.py --env_name CartPole-v0 -n 100 -b 5000 \
    -dsa --exp_name q1_lb_no_rtg_dsa

python rob831/scripts/run_hw2.py --env_name CartPole-v0 -n 100 -b 5000 \
    -rtg -dsa --exp_name q1_lb_rtg_dsa

python rob831/scripts/run_hw2.py --env_name CartPole-v0 -n 100 -b 5000 \
    -rtg --exp_name q1_lb_rtg_na
```
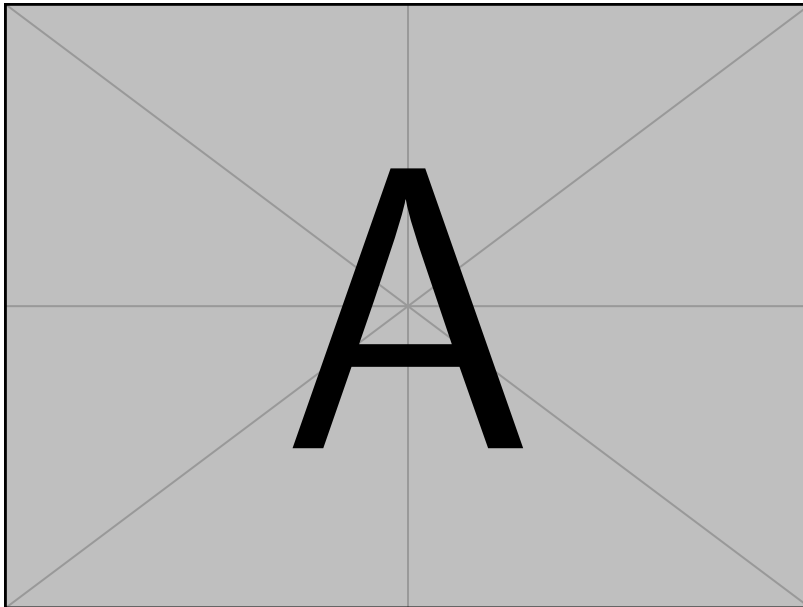---

5.1.2 Plots

5.1.2.1 Small batch – [5 points]

Q5.1.2.1

5.1.2.2   Large batch – [5 points]

Q5.1.2.2

5.1.3   Analysis

5.1.3.1   Value estimator – [5 points]

Q5.1.3.1

5.1.3.2   Advantage standardization – [5 points]

Q5.1.3.2

5.1.3.3   Batch size – [5 points]

> **Q5.1.3.1**
>
>

5.2   Experiment 2 (InvertedPendulum) – [15 points total]

5.2.1   Configurations – [5 points]

> **Q5.2.1**
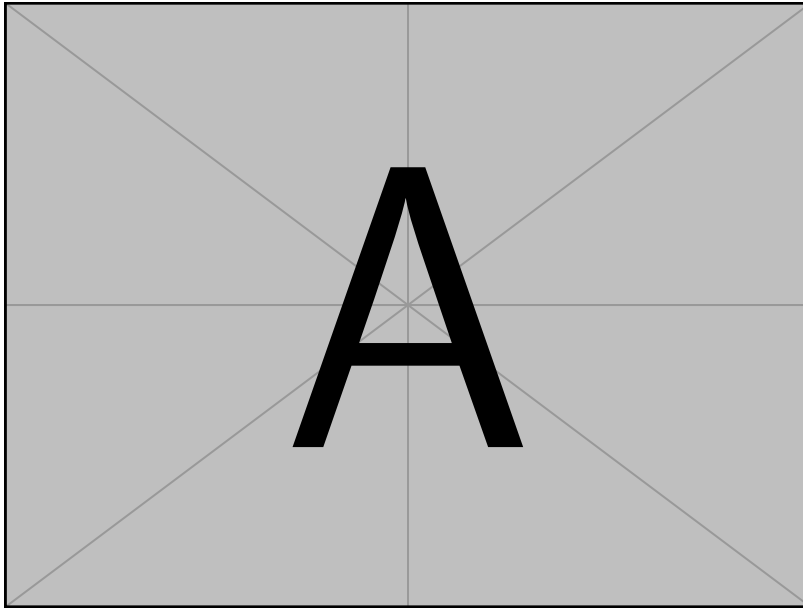>
> ```
> python rob831/scripts/run_hw2.py --env_name InvertedPendulum-v4 \
>     --ep_len 1000 --discount 0.9 -n 100 -l 2 -s 64 -b <b*> -lr <r*> -rtg \
>     --exp_name q2_b<b*>_r<r*>
> ```

5.2.2   smallest b* and largest r* (same run) – [5 points]

> **Q5.2.2**
>
>

5.2.3   Plot – [5 points]

Q5.2.3

7   More Complex Experiments

### 7.1   Experiment 3 (LunarLander) – [10 points total]

7.1.1   Configurations

Q7.1.1

```
python rob831/scripts/run_hw2.py \
    --env_name LunarLanderContinuous-v4 --ep_len 1000
    --discount 0.99 -n 100 -l 2 -s 64 -b 40000 -lr 0.005 \
    --reward_to_go --nn_baseline --exp_name q3_b40000_r0.005
```
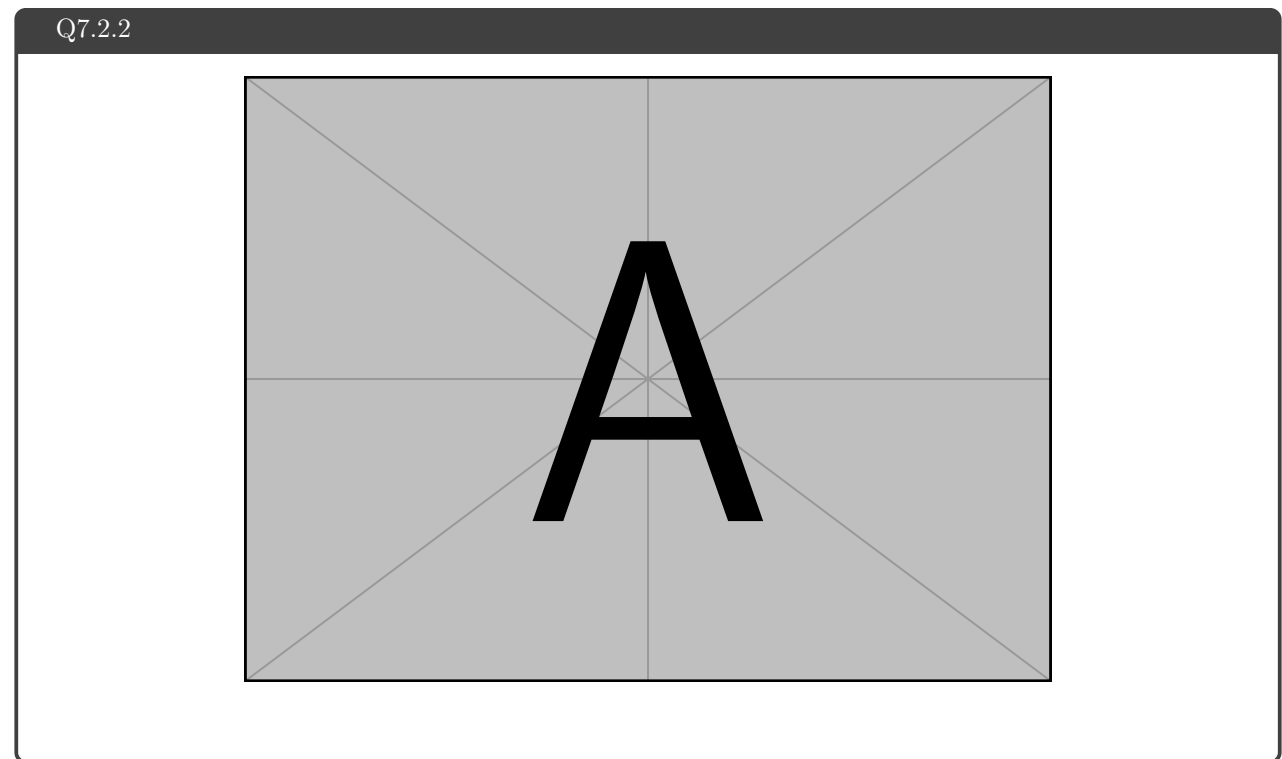
7.1.2   Plot – [10 points]

> **Q7.1.2**
>
>

7.2   Experiment 4 (HalfCheetah) – [30 points]

7.2.1   Configurations

> **Q7.2.1**
>
> ```
> # b ∈ [10000, 30000, 50000], r ∈ [0.005, 0.01, 0.02]
> python rob831/scripts/run_hw2.py --env_name HalfCheetah-v4 --ep_len 150 \
>     --discount 0.95 -n 100 -l 2 -s 32 -b <b> -lr <r> -rtg --nn_baseline \
>     --exp_name q4_search_b<b>_lr<r>_rtg_nnbaseline
> ```

7.2.2   Plot – [10 points]

Q7.2.2

7.2.3   Optimal b* and r* – [3 points]

Q7.2.3

7.2.4   Describe how b* and r* affect task performance – [7 points]

Q7.2.4

7.2.5   Configurations with optimal b* and r* – [3 points]

Q7.2.5

```
python rob831/scripts/run_hw2.py --env_name HalfCheetah-v4 --ep_len 150 \
    --discount 0.95 -n 100 -l 2 -s 32 -b <b*> -lr <r*> \
    --exp_name q4_b<b*>_r<r*>

python rob831/scripts/run_hw2.py --env_name HalfCheetah-v4 --ep_len 150 \
    --discount 0.95 -n 100 -l 2 -s 32 -b <b*> -lr <r*> -rtg \
    --exp_name q4_b<b*>_r<r*>_rtg

python rob831/scripts/run_hw2.py --env_name HalfCheetah-v4 --ep_len 150 \
    --discount 0.95 -n 100 -l 2 -s 32 -b <b*> -lr <r*> --nn_baseline \
    --exp_name q4_b<b*>_r<r*>_nnbaseline

python rob831/scripts/run_hw2.py --env_name HalfCheetah-v4 --ep_len 150 \
    --discount 0.95 -n 100 -l 2 -s 32 -b <b*> -lr <r*> -rtg --nn_baseline \
    --exp_name q4_b<b*>_r<r*>_rtg_nnbaseline
```

7.2.6   Plot for four runs with optimal b* and r* – [7 points]

Q7.2.6

8   Implementing Generalized Advantage Estimation

8.1 Experiment 5 (Hopper) – [20 points]

8.1.1 Configurations

Q8.1.1

```
# λ ∈ [0, 0.95, 0.99, 1]
python rob831/scripts/run_hw2.py \
    --env_name Hopper-v4 --ep_len 1000
    --discount 0.99 -n 300 -l 2 -s 32 -b 2000 -lr 0.001 \
    --reward_to_go --nn_baseline --action_noise_std 0.5 --gae_lambda <λ> \
    --exp_name q5_b2000_r0.001_lambda<λ>
```

8.1.2 Plot – [13 points]

Q8.1.2

8.1.3 Describe how λ affects task performance – [7 points]

Q8.1.3

9   Bonus! (optional)

9.1   Parallelization – [15 points]


Q9.1

Difference in training time:

```
python rob831/scripts/run_hw2.py \
```

9.2   Multiple gradient steps – [5 points]

Q9.1


```
python rob831/scripts/run_hw2.py \
```